

16-782

Planning & Decision-making in Robotics

***Interleaving Planning & Execution:
Real-time Heuristic Search***

Maxim Likhachev

Robotics Institute

Carnegie Mellon University

Planning during Execution

- Planning is a repeated process!
 - partially-known environments
 - dynamic environments
 - imperfect execution of plans
 - imprecise localization
- Need to be able to re-plan fast!
- Several methodologies to achieve this:
 - anytime heuristic search: return the best plan possible within T msecs
 - incremental heuristic search: speed up search by reusing previous efforts
 - **real-time heuristic search: plan few steps towards the goal and re-plan later**



this lecture


Real-time (Agent-centered) Heuristic Search


Enforce a strict limit on the amount of computations (no requirement on planning all the way to the goal)

Real-time (Agent-centered) Heuristic Search

1. Compute a partial path by expanding at most N states around the robot
2. Move once, incorporate sensor information, and goto step 1

Example in a fully-known terrain:







 - *expanded*

Real-time (Agent-centered) Heuristic Search

1. Compute a partial path by expanding at most N states around the robot
2. Move once, incorporate sensor information, and goto step 1

Example in an unknown terrain (planning with Freespace Assumption):





 - *expanded*

Planning with Freespace Assumption [Nourbakhsh & Genesereth, '96]

- Freespace Assumption: all unknown cells are assumed to be traversable
- Planning with the Freespace Assumption: always move the robot on a shortest path to the goal assuming all unknown cells are traversable
- Replan the path whenever a new sensor information received


*costs between unknown states is
the same as the costs in between
states known to be free*



 - unknown

Planning with Freespace Assumption [Nourbakhsh & Genesereth, '96]

- Freespace Assumption: all unknown cells are assumed to be traversable
- Planning with the Freespace Assumption: always move the robot on a shortest path to the goal assuming all unknown cells are traversable
- Replan the path whenever a new sensor information received



Real-time (Agent-centered) Heuristic Search

1. Compute a partial path by expanding at most N states around the robot
2. Move once, incorporate sensor information, and goto step 1

Research issues:

- how to compute partial path
- how to guarantee complete behavior (guarantee to reach the goal)
- provide bounds on the number of steps before reaching the goal

Real-time (Agent-centered) Heuristic Search

1. Compute a partial path by expanding at most N states around the robot
2. Move once, incorporate sensor information, and goto step 1

Research issues:

- how to compute partial path *Any ideas?*
- how to guarantee complete behavior (guarantee to reach the goal)
- provide bounds on the number of steps before reaching the goal

Learning Real-Time A* (LRTA*) [Korf, '90]

- Repeatedly move the robot to the most promising adjacent state, using heuristics

1. *always move as follows: $s_{start} = \operatorname{argmin}_{s \in \operatorname{succ}(s_{start})} c(s_{start}, s) + h(s)$*

$$h(x,y) = \max(\operatorname{abs}(x-x_{goal}), \operatorname{abs}(y-y_{goal})) + 0.4 * \min(\operatorname{abs}(x-x_{goal}), \operatorname{abs}(y-y_{goal}))$$

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4	3.4	2.4	1.4	1
5	4	3	2	1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4		2.4	1.4	1
5	4	3	2	1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	4	3		1	0

Any problems?

Learning Real-Time A* (LRTA*) [Korf, '90]

- Repeatedly move the robot to the most promising adjacent state, using heuristics

1. *always move as follows: $s_{start} = \operatorname{argmin}_{s \in \operatorname{succ}(s_{start})} c(s_{start}, s) + h(s)$*

$$h(x,y) = \max(\operatorname{abs}(x-x_{goal}), \operatorname{abs}(y-y_{goal})) + 0.4 * \min(\operatorname{abs}(x-x_{goal}), \operatorname{abs}(y-y_{goal}))$$

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	4	3		1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	4	3		1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	4	3		1	0

...

Local minima problem (myopic or incomplete behavior)

Any solutions?

Learning Real-Time A* (LRTA*) [Korf, '90]

- Repeatedly move the robot to the most promising adjacent state, using **and updating** heuristics

makes h-values more informed

1. *update* $h(s_{start}) = \min_{s \in \text{succ}(s_{start})} c(s_{start}, s) + h(s)$
2. *always move as follows:* $s_{start} = \text{argmin}_{s \in \text{succ}(s_{start})} c(s_{start}, s) + h(s)$

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4	3.4	2.4	1.4	1
5	4	3	2	1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4		2.4	1.4	1
5	4	3	2	1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	4	5		1	0

Learning Real-Time A* (LRTA*) [Korf, '90]

- Repeatedly move the robot to the most promising adjacent state, using **and updating** heuristics

1. *update* $h(s_{start}) = \min_{s \in \text{succ}(s_{start})} c(s_{start}, s) + h(s)$
2. *always move as follows*: $s_{start} = \text{argmin}_{s \in \text{succ}(s_{start})} c(s_{start}, s) + h(s)$

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	5.4	5		1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	5.2			1.4	1
5	5.4	5		1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	5.4	5		1	0

...

Learning Real-Time A* (LRTA*) [Korf, '90]

- Repeatedly move the robot to the most promising adjacent state, using **and updating** heuristics

1. *update* $h(s_{start}) = \min_{s \in succ(s_{start})} c(s_{start}, s) + h(s)$
2. *always move as follows:* $s_{start} = \operatorname{argmin}_{s \in succ(s_{start})} c(s_{start}, s) + h(s)$

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	5.4	5		1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	5.2			1.4	1
5	5.4	5		1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	5.4	5		1	0

...

h -values remain admissible and consistent

proof?

Learning Real-Time A* (LRTA*) [Korf, '90]

- Repeatedly move the robot to the most promising adjacent state, using **and updating** heuristics

1. *update* $h(s_{start}) = \min_{s \in \text{succ}(s_{start})} c(s_{start}, s) + h(s)$
2. *always move as follows*: $s_{start} = \text{argmin}_{s \in \text{succ}(s_{start})} c(s_{start}, s) + h(s)$

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	5.4	5		1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	5.2			1.4	1
5	5.4	5		1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	5.4	5		1	0

...

robot is guaranteed to reach goal in finite number of steps if:

- all costs are bounded from below with $\Delta > 0$
- graph is of finite size and there exists a finite-cost path to the goal
- all actions are reversible

Learning Real-Time A* (LRTA*) [Korf, '90]

- Repeatedly move the robot to the most promising adjacent state, using **and updating** heuristics

1. *update* $h(s_{start}) = \min_{s \in \text{succ}(s_{start})} c(s_{start}, s) + h(s)$
2. *always move as follows:* $s_{start} = \text{argmin}_{s \in \text{succ}(s_{start})} c(s_{start}, s) + h(s)$

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	5.4	5		1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	5.2			1.4	1
5	5.4	5		1	0

6.2	5.2	4.2	3.8	3.4	3
5.8	4.8	3.8	2.8	2.4	2
5.4	4.4			1.4	1
5	5.4	5		1	0

...

robot is guaranteed to reach goal in finite number of steps if:

- all costs are bounded from below with $\Delta > 0$
- graph is of finite size and there exists a finite-cost path to the goal
- all actions are reversible

Why conditions?

Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands [Koenig, '04]
 - 1. expand N states
 - 2. update h -values of expanded states by Dynamic Programming (DP)
 - 3. move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$

*necessary for the guarantee
to reach the goal*



- *expanded*

Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

state s :

- the state that minimizes cost to it plus heuristic estimate of the remaining distance
- the state that looks most promising in terms of the whole path from current robot state to goal



- *expanded*

Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
5	4		2	1
4	3	2		0

4-connected grid (robot moves in 4 directions)

example borrowed from ICAPS'06 planning summer school lecture (Koenig & Likhachev)



- *expanded*


Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6			3	2
			2	1
				0


expand $N=7$ states

 - *expanded*

Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*




Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands


How path is found?

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*



expand $N=7$ states

unexpanded state with smallest $g + h (= 5 + 3)$

 - *expanded*


Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	∞	∞	3	2
∞	∞		2	1
∞	∞	∞		0

*update h -values of expanded states via DP:
compute $h(s) = \min_{s' \in \text{succ}(s)} (c(s, s') + h(s'))$
until convergence*

 - *expanded*

Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	∞	4	3	2
∞	∞		2	1
∞	∞	∞		0

*update h -values of expanded states via DP:
compute $h(s) = \min_{s' \in \text{succ}(s)} (c(s, s') + h(s'))$
until convergence*



- *expanded*


Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
∞	∞		2	1
∞	∞	∞		0

*update h -values of expanded states via DP:
compute $h(s) = \min_{s' \in \text{succ}(s)} (c(s, s') + h(s'))$
until convergence*

 - *expanded*


Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
∞	6		2	1
∞	∞	∞		0

*update h -values of expanded states via DP:
compute $h(s) = \min_{s' \in \text{succ}(s)} (c(s, s') + h(s'))$
until convergence*

 - *expanded*

Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
7	6		2	1
∞	∞	∞		0

*update h -values of expanded states via DP:
compute $h(s) = \min_{s' \in \text{succ}(s)} (c(s, s') + h(s'))$
until convergence*

*Does it matter in
what order?*



- expanded

Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
7	6		2	1
∞	7	∞		0

*update h -values of expanded states via DP:
compute $h(s) = \min_{s' \in \text{succ}(s)} (c(s, s') + h(s'))$
until convergence*



- *expanded*


Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
7	6		2	1
8	7	∞		0

*update h -values of expanded states via DP:
compute $h(s) = \min_{s' \in \text{succ}(s)} (c(s, s') + h(s'))$
until convergence*

 - *expanded*


Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
7	6		2	1
8	7	8		0

*update h -values of expanded states via DP:
compute $h(s) = \min_{s' \in \text{succ}(s)} (c(s, s') + h(s'))$
until convergence*

 - *expanded*

Learning Real-Time A* (LRTA*)

- LRTA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states by Dynamic Programming (DP)*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
7	6		2	1
8	7	8		0

*make a move along the found path
and repeat steps 1-3*

*Drawbacks compared
to A*?*




- expanded

Real-time Adaptive A* (RTAA*) [Koenig & Likhachev, '06]

- RTAA* with $N \geq 1$ expands


1. expand N states
2. update h -values of expanded states u by $h(u) = f(s) - g(u)$,
where $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$
3. move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$

*one linear pass,
and even that can be lazy(postponed)*



expand $N=7$ states

unexpanded state s with smallest
 $g + h (= 5 + 3)$

 - expanded

Real-time Adaptive A* (RTAA*)

- RTAA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states u by $h(u) = f(s) - g(u)$,
where $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	g=3	g=4	3	2
g=3	g=2		2	1
g=2	g=1	g=0		0

update all expanded states u :
 $h(u) = f(s) - g(u)$

unexpanded state s with smallest
 $f(s) = 8$



- *expanded*

Real-time Adaptive A* (RTAA*)

- RTAA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states u by $h(u) = f(s) - g(u)$,
where $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	8-3	8-4	3	2
8-3	8-2		2	1
8-2	8-1	8-0		0

update all expanded states u :
 $h(u) = f(s) - g(u)$

unexpanded state s with smallest
 $f(s) = 8$



- *expanded*

Real-time Adaptive A* (RTAA*)

- RTAA* with $N \geq 1$ expands

1. *expand N states*
2. *update h -values of expanded states u by $h(u) = f(s) - g(u)$,
where $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*
3. *move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
5	6		2	1
6	7	8		0

update all expanded states u :
 $h(u) = f(s) - g(u)$

unexpanded state s with smallest
 $f(s) = 8$



- *expanded*

Real-time Adaptive A* (RTAA*)

- RTAA* with $N \geq 1$ expands

1. expand N states
2. update h -values of expanded states u by $h(u) = f(s) - g(u)$,
where $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$
3. move on the path to state $s = \operatorname{argmin}_{s' \in \text{OPEN}} g(s') + h(s')$

proof of admissibility:

$$g(u) + h^*(u) \geq h^*(s_{\text{start}})$$

$$h^*(u) \geq h^*(s_{\text{start}}) - g(u)$$


$$h^*(u) \geq f(s) - g(u)$$

$$h^*(u) \geq h_{\text{updated}}(u)$$

$h^*(\cdot)$ – true cost-to-goal

because $f(s) \leq h^*(s_{\text{start}})$

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
5	6		2	1
6	7	8		0

 - expanded

LRTA* vs. RTAA*

LRTA*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
7	6		2	1
8	7	8		0

RTAA*

8	7	6	5	4
7	6	5	4	3
6	5	4	3	2
5	6		2	1
6	7	8		0

- Update of h -values in RTAA* is much faster but not as informed
- Both guarantee admissibility and consistency of heuristics
- For both, heuristics are monotonically increasing
- Both guarantee to reach the goal in a finite number of steps (given the conditions listed previously)

What You Should Know...

- What Freespace Assumption means
- Why we need to update heuristics in the context of Real-time Heuristic Search
- The operation of LRTA*
- Pros and cons of LRTA*
- What domains LRTA* is useful in and what domains it is not really applicable
- What RTAA* is